

**A REPORT  
ON  
VISUAL QUESTION ANSWERING**

**BY  
AASHISH AGRAWAL (2015B3A80411G)  
ADIT VINAY DESHMUKH (2015B3A70400G)  
VARUN B YELIGAR (2015B3A70403G)**

**Prepared in partial fulfillment of the  
Neural Networks and Fuzzy Logic  
(BITS F312)**

**AT**



**BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE, GOA CAMPUS  
NOVEMBER, 2018**

## Basic Model

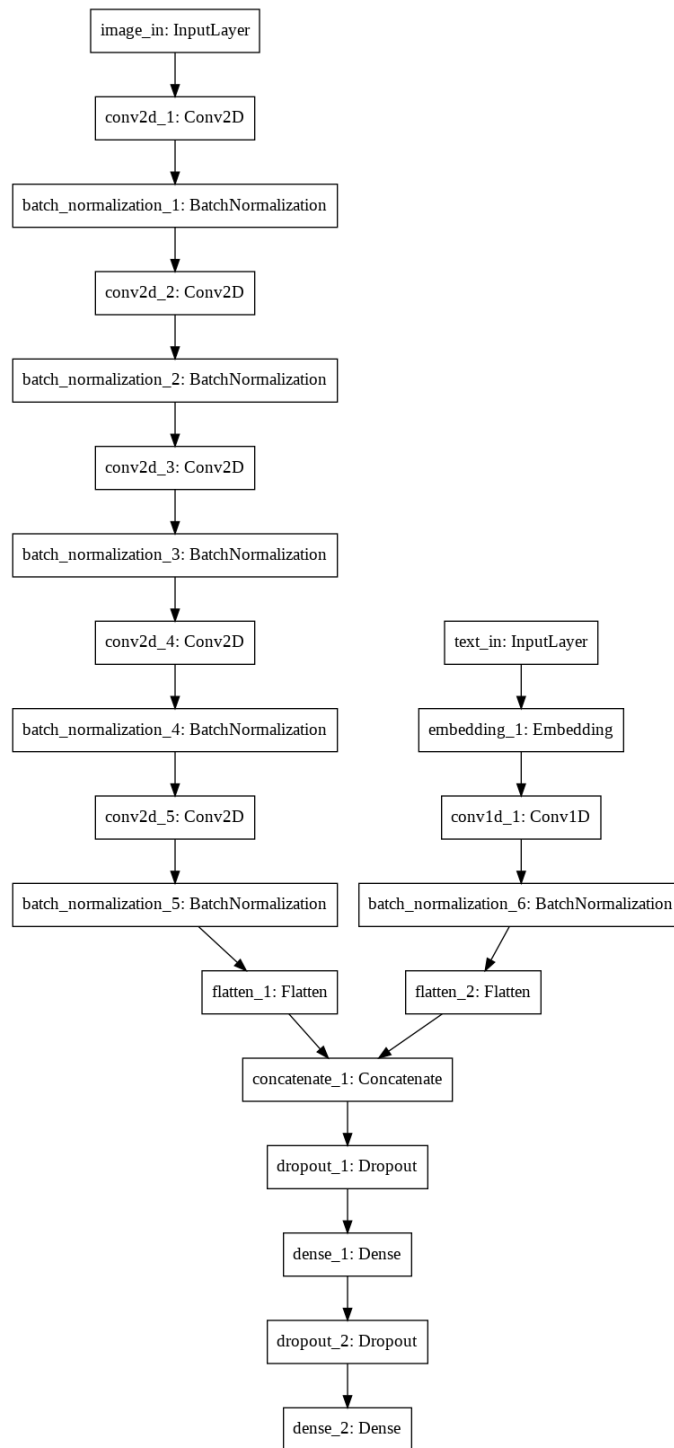
Initially to understand the difficulty and challenges we may face in the task, we designed the most basic of a model which may solve the problem. Since both the image size and the text sequence length were small to begin with, no preprocessing was done on either of them, to get as much information out of them as possible. For constructing the model, we used simple Conv2D layers one after the other which were later flattened, for the purpose of encoding images. Conv1D layers were used to encode the questions text data and was later flattened.

After the encoding of both the text and images data, the encoded vectors thus obtained were concatenated to form a single dense layer. One more Dense layer was appended on top of it and dropouts were used in between them for regularization purposes.

The model summary of the data is presented below:

Layer (type)	Output Shape	Param #	Connected to
image_in (InputLayer)	(None, 120, 160, 3)	0	
conv2d_1 (Conv2D)	(None, 59, 79, 24)	672	image_in[0][0]
batch_normalization_1 (BatchNor	(None, 59, 79, 24)	96	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, 29, 39, 48)	10416	batch_normalization_1[0][0]
batch_normalization_2 (BatchNor	(None, 29, 39, 48)	192	conv2d_2[0][0]
conv2d_3 (Conv2D)	(None, 14, 19, 48)	20784	batch_normalization_2[0][0]
batch_normalization_3 (BatchNor	(None, 14, 19, 48)	192	conv2d_3[0][0]
text_in (InputLayer)	(None, 48)	0	
conv2d_4 (Conv2D)	(None, 6, 9, 64)	27712	batch_normalization_3[0][0]
embedding_1 (Embedding)	(None, 48, 256)	262144	text_in[0][0]
batch_normalization_4 (BatchNor	(None, 6, 9, 64)	256	conv2d_4[0][0]
conv1d_1 (Conv1D)	(None, 23, 128)	98432	embedding_1[0][0]
conv2d_5 (Conv2D)	(None, 2, 4, 64)	36928	batch_normalization_4[0][0]
batch_normalization_6 (BatchNor	(None, 23, 128)	512	conv1d_1[0][0]
batch_normalization_5 (BatchNor	(None, 2, 4, 64)	256	conv2d_5[0][0]
flatten_2 (Flatten)	(None, 2944)	0	batch_normalization_6[0][0]
flatten_1 (Flatten)	(None, 512)	0	batch_normalization_5[0][0]
concatenate_1 (Concatenate)	(None, 3456)	0	flatten_2[0][0] flatten_1[0][0]
dropout_1 (Dropout)	(None, 3456)	0	concatenate_1[0][0]
dense_1 (Dense)	(None, 256)	884992	dropout_1[0][0]
dropout_2 (Dropout)	(None, 256)	0	dense_1[0][0]
dense_2 (Dense)	(None, 26)	6682	dropout_2[0][0]
Total params: 1,350,266			
Trainable params: 1,349,514			
Non-trainable params: 752			

The flow chart of the model is presented below:



Model thus compiled was trained for 50 epochs, and the resultant files (tokenizer.pickle and model.h5) were stored.

The following graphs plot, accuracy vs epochs and loss vs epochs for both training and validation data. The graph, however, is plotted only for the first 10 epochs, connectivity issues on Google Colab.



Before we ended up this model, LSTMs were used instead of simple Con1D layers, to encode the text data, however that model was rejected in the initial training phases because of lengthy amount of time taken by it during training and bigger model size, combined with below average results. So, the CNN+LSTM model was discarded in the initial phases only.


## Relational Network approach:

The current state of the art result on the CLEVR dataset, of which we have been given a subset, was achieved using relational networks proposed by scientists at Deep Mind. They claimed to have achieved an overall accuracy of 95.5% beating the previous best of 77%. They proposed a new architecture which was built to understand relational dependencies just as convolutional networks are built to understand spatial properties and RNNs are built to understand sequential dependencies. One of the main advantages of the relational model is that it is much smaller in size compared to the previous models and hence is computationally efficient.

The model that we built based on the paper was only 1.8MB in size and achieved accuracy of 45.44% on the validation set. The relational model is smaller but performs better than the other models because it is built to find relational relations and because the complexity in the task is not that the images are complicated but that finding the relations is hard. The relational model addresses this difficulty.

We found that a simpler model which did not use relational networks worked better in practice on the dataset that we had been provided.

The model described in the paper:



Relational model mathematically:


$$\text{RN}(O) = f_{\phi} \left( \sum_{i,j} g_{\theta}(o_i, o_j) \right)$$

The summary of the model made in PyTorch:

```
DataParallel(  
(module): RelationNetworks(  
  (conv): Sequential(  
    (0): Conv2d(3, 24, kernel_size=[3, 3], stride=(2, 2), padding=(1, 1), bias=False)  
    (1): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
    (3): Conv2d(24, 24, kernel_size=[3, 3], stride=(2, 2), padding=(1, 1), bias=False)  
    (4): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU()  
    (6): Conv2d(24, 24, kernel_size=[3, 3], stride=(2, 2), padding=(1, 1), bias=False)  
    (7): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (8): ReLU()  
    (9): Conv2d(24, 24, kernel_size=[3, 3], stride=(2, 2), padding=(1, 1), bias=False)  
    (10): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (11): ReLU()  
  )  
  (embed): Embedding(90, 32)  
  (lstm): LSTM(32, 128, batch_first=True)  
  (g): Sequential(  
    (0): Linear(in_features=180, out_features=256, bias=True)  
    (1): ReLU()  
    (2): Linear(in_features=256, out_features=256, bias=True)  
    (3): ReLU()  
    (4): Linear(in_features=256, out_features=256, bias=True)  
    (5): ReLU()  
    (6): Linear(in_features=256, out_features=256, bias=True)  
    (7): ReLU()  
  )  
  (f): Sequential(  

```

(0): Linear(in\_features=256, out\_features=256, bias=True)  
 (1): ReLU()  
 (2): Linear(in\_features=256, out\_features=256, bias=True)  
 (3): ReLU()  
 (4): Dropout(p=0.5)  
 (5): Linear(in\_features=256, out\_features=29, bias=True)  
 )  
 )  
 )



$$\text{RN}(O) = f_{\phi} \left( \sum_{i,j} g_{\theta}(o_i, o_j) \right)$$

The current state of the art result on the CLEVR dataset, of which we have been given a subset, was achieved using relational networks proposed by scientists at Deep Mind. They claimed to have achieved an overall accuracy of 95.5% beating the previous best of 77%. They proposed a new architecture which was built to understand relational dependencies just as convolutional networks are built to understand spatial properties and RNNs are built to understand sequential dependencies. One of the main advantages of the relational model is that it is much smaller in size compared to the previous models and hence is computationally efficient.

The model that we built based on the paper was only 1.8MB in size and achieved accuracy of 45.44% on the validation set. The relational model is smaller but performs better than the other models because it is built to find relational relations and because the complexity in the task is not that the images are complicated but that finding the relations is hard. The relational model addresses this difficulty.

We found that a simpler model which did not use relational networks worked better in practice on the dataset that we had been provided.